

Q&A cheat sheet

Rami's section · Fairness, deployment, and limitations

Ten likely questions on your part of the talk, with answers tuned to ~20 seconds spoken. Read tonight, skim before the session. If a question isn't on this sheet, fall back to: **"No single model is best across communities, so deployment depends on cost and use case."**

1. How did you choose which communities to report?

"We only report groups with at least **15 test examples**, so per-group statistics are at least marginally interpretable. Three communities clear that bar: Blacks at 46, Others at 39, Jews at 15. Smaller groups exist in HateMM but we don't have the test support to characterise them."

Tip: If pushed about the missing groups, name the threshold: $n \geq 15$.

2. Is "Others" even a meaningful category?

"It's a **catch-all**, canonicalised from multi-target rows by alphabetic comma-join, so it mixes intersectional hate cases with edge labels. We report it for transparency, but we treat its results as descriptive only, not conclusive. The poster notes this explicitly."

Tip: Use the word "descriptive" to signal you're not over-claiming.

3. With only 15 Jews-targeted examples, isn't the per-group conclusion unreliable?

"Yes, and we say so on the poster. The seed standard deviation for fusion on that group is plus-minus 0.13, an **order of magnitude larger** than its mean advantage. Any claim at $n = 15$ is descriptive, not statistical."

Tip: Don't defend the number. Acknowledge the limitation and move on.

4. Why does text alone win on Jews and Others while fusion wins on Blacks?

"Hate against Jews and Others in this dataset tends to be **more verbally explicit**, so the transcript captures most of the discriminating signal. For Blacks, the signal seems more spread across modalities, visual symbols, audio cues, so fusion's ranking gain converts to a threshold-level win. We don't claim this pattern generalises beyond HateMM."

Tip: Anchor in the data, not in a sociological claim.

5. Why does 85% accuracy on "Others" indicate the model is defaulting, not learning?

"Because that subset is **heavily skewed toward non-hate**, the model can hit 85% accuracy just by defaulting to 'non-hate' on every example. Macro-F1, which weights hate and non-hate equally, collapses to about 0.5 on this group despite the high accuracy. The accuracy is real but uninformative."

Tip: Punchline: "high accuracy with low F1 = model defaulting to majority class."

6. How should one of these models actually be deployed?

"Our recommendation is **triage, not autonomous moderation**. Use the AUROC-strong fused model to rank videos for a human moderator. The model improves ordering by risk; the human still makes the final yes/no call. That matches the modality where fusion actually wins, which is ranking quality."

Tip: "Triage, not autonomous moderation" is the line to remember.

7. Why is V+T the cost-aware sweet spot rather than full fusion?

"V+T reaches 0.86 AUROC, within 0.02 of full fusion at 0.88, which is roughly 60% of the AUROC lift over text-only. That capture rate comes **without needing an audio pipeline at all**: no ffmpeg extraction, no wav2vec2 inference, no Whisper transcription. For platforms without audio infrastructure, V+T is the practical choice."

Tip: Stress the deployment cost saving, not the marginal AUROC gain.

8. Why are the encoders frozen? Wouldn't fine-tuning do better?

"Fine-tuning would lift absolute numbers but **break the ablation**. Our goal was to isolate modality contribution, not to maximise raw performance. If a downstream user only cares about deployed accuracy, they would unfreeze. We chose interpretability over the last few points."

Tip: "Interpretability over the last few points" is the one-liner.

9. Why don't your results transfer to YouTube or TikTok?

"BitChute is a **low-moderation platform** with a hate base rate around 40%. Mainstream platforms run at 1 to 5%. A classifier calibrated on our distribution would be far too aggressive on a mainstream feed; we'd need to recalibrate the threshold and probably retrain on platform-specific data."

Tip: Quote the base rates: 40% on BitChute, 1 to 5% mainstream.

10. What would you change with two more weeks?

"Three things. First, **bootstrap CIs** on the test set to back the AUROC claim with a formal significance test. Second, more seeds plus optional five-fold cross-validation to tighten per-community estimates. Third, a **counterfactual modality-masking probe** to give moderators per-video modality attribution rather than only aggregate contributions."

Tip: Forward-looking answers signal research maturity. Land on the third item.

If you blank

- **Three lines that fit any question:** "No single model is best across communities." / "Use it for triage, not autonomous moderation." / "BitChute is a fringe distribution; results don't transfer unchanged."
- **If you don't know:** "Good question, that's not in our current analysis, but my intuition would be..." Honesty beats bluffing.
- **If two people ask at once:** Pick one, answer cleanly, then "and to your question..." Don't try to braid them.