

Poster rehearsal script

3 minutes total · 3 speakers · 60 seconds each

Order matches the left-to-right reading flow of the poster. You can swap roles, but keep this order so listeners' eyes track you naturally.

Speaker 1: problem, dataset, method. Speaker 2: results and the Seed Trap. Speaker 3: fairness, practical implication, limitations.

Speaker 1 · Youssef

approximately 60 s

Stand to the left of the poster. Point at the title.

"Watch. Listen. Read. Three things you do at the same time when you watch a video, and three things most hate-detection models still do separately. Our project asks: if a video can carry hate in any of those channels at once, can a single model handle all three together?"

Point at the Dataset block.

"We use **HateMM**, a 1,083-video benchmark from BitChute released by Das et al. in 2023. About forty percent are labeled hate, and every hate video carries a target-community tag, which lets us audit fairness later."

Move hand to the Method architecture diagram.

"Our pipeline is simple. Every video gets split into three streams: frames, audio, transcript. Each stream goes through a **frozen pretrained encoder**: CLIP for images, wav2vec2 for audio, RoBERTa for text. The only thing we actually train is the part in red, a small two-layer transformer that fuses the three streams and outputs hate or non-hate."

Step back. Hand off to Kevin.

Step to the right side of the poster. Point at Table 1.

"Here is what happens when we run all seven possible modality combinations. **AUROC, the right column, rises every single time we add a modality.** Audio alone, point seven seven. Video, point eight one. Text, point eight three. All three together, point eight eight. Each step buys ranking signal."

Point at the Key Observations callout.

"But look at accuracy. Text alone gets seventy-nine percent. Full fusion gets seventy-nine percent. **They tie.** So fusion only improves how the model **ranks videos by risk.** By 'risk' we mean the model's hate-score: fusion pulls suspect videos closer to the top of the queue, but the same number of videos end up labeled hate overall."

Move hand to the Seed Trap section on the bottom-left.

"And here is the story we almost missed. Our first run used seed 42 only and showed fusion winning by **five point five accuracy points.** Across three seeds that gap vanished completely. A single-seed paper would have shipped a false claim. **Only AUROC survived multi-seed scrutiny.**"

Step back. Hand off to Rami.

IF ASKED · WHAT IS AUROC?

"AUROC asks one question: if I pick a random hate video and a random non-hate video, how often does the hate one get a higher hate-score from the model? **1.0 means always** (perfect ranking), **0.5 is a coin flip.** Unlike accuracy, it does not depend on where you draw the cutoff between hate and non-hate; it only cares whether the ordering is right. That is why fusion can lift AUROC without lifting accuracy."

Point at the Fairness Analysis table.

"We then broke results down by the target community of each hate video. **Text alone actually wins on Jews and Others. Fusion only wins on Blacks.** No single model is best for every group, and the smallest community has only fifteen test examples, so per-group conclusions need care."

Point at the Conclusion and Practical Implication block.

"That gives us three honest deployment options. **Full fusion** when ranking quality matters most, for example to prioritise a moderator's queue. **Video-plus-text** as a cost-aware sweet spot, within two AUROC points of full fusion with no audio pipeline at all. **Text alone** if you just need a binary call at a fixed threshold."

Point at the Limitations list.

"Caveats. One hundred and nine test videos is small. BitChute is a fringe platform. Whisper degrades on shouting and accents. Frozen encoders cap absolute performance. So take the shape of our results as the contribution, not the final numbers. **Thank you.**"

Pause. Smile. Wait for questions.

Delivery notes for all three of you

- **Rehearse the handoffs out loud at least three times.** Silent pauses kill momentum more than slightly fluffed words.
- **Do not read the poster.** Eyes on the listener for at least half of each sentence.
- **Numbers:** say "point eight eight" not "zero point eight eight". Faster, more natural.
- **Pointers:** actually move your hand to the element you are naming. Examiners track gestures.
- **If interrupted mid-script:** the speaker talking answers the question, then resumes from the next sentence. The other two stay quiet unless invited in.

AUROC up · accuracy flat · Seed Trap

If you blank, these four words let you reconstruct the whole talk.